

Adithyan K

AI Systems Engineer | LLM Orchestration & Agent Infrastructure
adithyn.krishna@gmail.com | +91 7034046267 | adithyn.dev | [GitHub](https://github.com/adithyn) | [LinkedIn](https://www.linkedin.com/in/adithyan-k)

EDUCATION

Manipal Institute of Technology

Bengaluru, India

B.Tech in Computer Science and Engineering (AI) – **CGPA: 8.36**

Sep. 2022 – Present

- **Relevant Coursework:** Machine Learning, Deep Learning, Natural Language Processing, Generative AI, Reinforcement Learning, Data Structures & Algorithms

EXPERIENCE

AI Engineer Intern

Jan 2026 – Present

SuperAGI – AI native platform for sales, marketing, and support automation

Bengaluru, India

- Built the execution and deployment infrastructure for **Vibe Coder**, SuperAGI's production AI app builder – integrated **Modal** sandboxes with build/serve/proxy tooling that runs LLM-generated full-stack apps and streams a live preview from a natural-language prompt; **450+ apps generated and served**.
- Shipped the static-app deployment pipeline on **S3 + CloudFront**, using **CloudFront Functions** for per-app subdomain routing at the edge, and designed the agent's execution environment – the project scaffold the code-generation agent operates within.
- Built the backend for **Digital Employees**, autonomous voice + avatar agents that join client meetings to run sales demos – bridged the agent loop into **LiveKit**'s real-time voice pipeline (managing the STT WebSocket) and drove segment-wise slide transitions via LLM structured outputs; used across **4,000+ meetings** in production.

Co-Founder & Lead AI Engineer

Jul 2025 – Present

Lifie – Voice and chat AI for businesses

Remote

- Co-founded a voice + chat AI platform for businesses – **REACH** (outbound voice engagement) and **ASSIST** (inbound chat support and sales workflows).
- Built the entire AI + backend layer of **ASSIST** – a multi-tenant **Mastra**/TypeScript agent platform, live with businesses in production, combining per-business RAG, tool-calling, and configurable booking/sales/support workflows delivered over the direct **WhatsApp Cloud API**.
- Engineered retrieval grounding with **Qdrant + FastEmbed** and multi-pass LLM guardrail validation; built **Lifie Hub**, an agent-discovery layer with vector search and **A2A + MCP** for cross-platform agent interoperability.

AI Product Development Intern

June 2025 – July 2025

QpiAI

Bengaluru, India

- Built and shipped a complete sales support agent on QpiAI's Agent Hive platform – automated lead research, enrichment, and personalized outreach generation, replacing a previously manual workflow.
- Integrated LLM APIs and internal SLMs for dynamic lead scoring, enrichment, and content generation across the pipeline.

PROJECTS

Metafog.io | *Next.js, Docker, Nginx, Ollama, LlamaIndex*

[Live Link](#)

- Built the full stack of a decentralized AI search engine – a crawl → retrieve → LLM-summarize (LlamaIndex) pipeline served by a network of **external provider nodes** contributing GPU inference (Ollama) and bandwidth; owned frontend, backend, Docker orchestration, and Nginx load balancing across provider nodes.

LLM Fine-tuning – ArXiv Research Q&A | *Unsloth, QLoRA, Hugging Face, LangChain, Qdrant, Docker GitHub*

- Fine-tuned **Qwen 8B & 14B** with **QLoRA + Unsloth** on a domain-specific research corpus and built the end-to-end pipeline – PDF ingestion, chunking, embedding, and **Qdrant + LangChain** retrieval-grounded inference to reduce out-of-context hallucinations.

SKILLS

Core: AI Systems Engineering, LLM Orchestration, Agent Infrastructure, RAG Pipelines, Backend Architecture

Languages: Python, TypeScript, Go, JavaScript, SQL

AI/ML: LLM Fine-tuning (QLoRA, Unsloth), Multi-Agent Systems, Mastra, LangChain, Context Engineering, Structured Outputs, MCP, A2A, Voice AI (LiveKit), Vector Search (Qdrant, FastEmbed)

Infrastructure: Docker, Modal, AWS (S3, CloudFront), GCP, Nginx, FastAPI, GitHub Actions, PostgreSQL, Qdrant, ChromaDB